




We can think about a **Set** in 2 different ways.

(1)



Graph without any edges

(2)



Graph that is fully connected

Message Passing
: Allow every node talk to every other node
: This is the core principle behind transformer networks

Quick Recap

- Within linear algebra, each permutation defines a $|V| \times |V|$ matrix.
- Such matrices are called permutation matrices.
- They have exactly one 1 in every row and column, zeroes elsewhere

Ex)

$$P_{(2,4,1,3)} X = \begin{bmatrix} 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} -x_1 \\ -x_2 \\ -x_3 \\ -x_4 \end{bmatrix} = \begin{bmatrix} -x_2 \\ -x_4 \\ -x_1 \\ -x_3 \end{bmatrix}$$

Learning on Sets

① Set-up

For now, assume our graph has no edges (i.e., $E = \emptyset$, the set of nodes)

Let $x_i \in \mathbb{R}^k$ be the features of node i . (i.e., our feature space is $C = \mathbb{R}^k$)

We can stack these features into a node feature matrix of shape $N \times k$. $X = (x_1, \dots, x_N)^T$

• i^{th} row of X corresponds to x_i

Permutation invariant operator is an operator (F), that if we apply F to our set or a permuted version of a set, we're getting the same output.

$$f\left(\begin{pmatrix} x_1 \\ x_5 \\ x_4 \\ x_3 \\ x_2 \end{pmatrix}\right) = y = f\left(\begin{pmatrix} x_5 \\ x_1 \\ x_4 \\ x_3 \\ x_2 \end{pmatrix}\right)$$

Symmetry group G : n -element permutation group S_n

Group element $g \in G$: permutation

Permutation invariance

Want: function $f(X)$ over sets that will not depend on the order

Equivalently: Applying a permutation matrix shouldn't modify result

$f(X)$ is permutation invariant if, for all permutation matrices P : $f(PX) = f(X)$

Permutation Equivariance:

• Permutation invariant models are good for **set-level outputs**. If we would like to **answer at the node level**, we need **permutation equivariant models**.

• Permutation equivariant functions: $F(PX) = PF(X)$

↳ It doesn't matter if we do permute before F or later F .

Deep Sets

$$f(X) = \phi\left(\bigoplus_{i \in V} \psi(x_i)\right)$$


where ϕ, ψ are (learnable) functions, e.g. MLPs

\bigoplus denotes any permutation-invariant operator.

General blueprint for learning on sets


One way we can enforce locality in equivariant set functions is through a shared function ψ applied to every node in isolation: $h_i = \psi(x_i)$, and stack h_i into a matrix $H = F(X)$

Pictorial View



Note: ψ : a shared function to every node in isolation
And, there's no message passing

Now, (2): Sets as fully connected graphs



$$\text{Message Passing: } f(x_i) = \phi\left(x_i, \bigoplus_{j \in N_i} \psi(x_i, x_j)\right)$$

$\psi(x_i, x_j)$: Message Passing between x_i and x_j

$\bigoplus_{j \in N_i}$: Message Aggregation

$\phi(x_i, z)$: Node feature update

Basic Self-Attention

Input: Sequence of tensors x_1, x_2, \dots, x_t

Output: Sequence of tensors, each one a weighted sum of the input sequence: y_1, y_2, \dots, y_t


$$y_i = \sum_j w_{ij} x_j$$

w in self-attention is not a learned weight, but a function of x_i and x_j : $w'_{ij} = x_i^T x_j$

w must sum to 1 over j

$$w_{ij} = \frac{\exp w'_{ij}}{\sum_j \exp w'_{ij}} : \text{Softmax (we're applying softmax to the nodes in some sequence)}$$

Pictorial View of basic self attention



$$y_2 = \sum_j w_{2j} x_j = w_{21}x_1 + w_{22}x_2 + w_{23}x_3 + w_{24}x_4$$

$$= \frac{\exp(x_2^T x_1)}{\sum_j \exp(x_2^T x_j)} x_1 + \frac{\exp(x_2^T x_2)}{\sum_j \exp(x_2^T x_j)} x_2 + \frac{\exp(x_2^T x_3)}{\sum_j \exp(x_2^T x_j)} x_3 + \frac{\exp(x_2^T x_4)}{\sum_j \exp(x_2^T x_j)} x_4 = 0.1x_1 + 0.8x_2 + 0 + 0.1x_4$$

★ $x_i^T x_j$: Notion of similarity between x_i and x_j

Therefore, w_{2j} tells how important x_j is to x_2 . And y_2 is computed based on its relationship with other nodes. This is "Attention"!

Basic Self Attention


No learned weights

Order of sequence doesn't affect result of computations

Query, Key, Value

: Every input vector x_i is used in 3 ways:

- Query: Compared to every other vector to compute attention weights for its own output y_i ;
- Key: Compared to every other vector to compute attention weight w_{ij} for output y_j ;
- Value: Summed with other vectors to form the result of the attention weighted sum





: We can process each input vector to fulfill the three roles with matrix multiplication

Learning the matrices \rightarrow learning attention

$$\begin{aligned} q_i &= W_q x_i & w'_{ij} &= q_i^T k_j \\ k_i &= W_k x_i & w_{ij} &= \text{Softmax}(w'_{ij}) \\ v_i &= W_v x_i & y_i &= \sum_j w_{ij} v_j \end{aligned}$$

W_q, W_k, W_v : Trainable matrices, such as MLP



$$\text{And, } w'_{ij} = q_i^T k_j$$

$$w_{ij} = \text{Softmax}(w'_{ij}) : \text{Scalar value.}$$


$$y_i = \sum_j w_{ij} v_j$$

$$W_q = \begin{bmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \end{bmatrix}$$

$$W_k = \begin{bmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \end{bmatrix}$$

$$\begin{aligned} q_i &= W_q x_i \\ &= \begin{bmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \end{bmatrix} \begin{bmatrix} x_1^i \\ x_2^i \\ x_3^i \end{bmatrix} \\ &= \begin{bmatrix} h_1 \\ h_2 \end{bmatrix}_q \end{aligned}$$

$$\begin{aligned} k_j &= W_k x_j \\ &= \begin{bmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \end{bmatrix} \begin{bmatrix} x_1^j \\ x_2^j \\ x_3^j \end{bmatrix} \\ &= \begin{bmatrix} h_1 \\ h_2 \end{bmatrix}_k \end{aligned}$$



$$W_v = \begin{bmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \end{bmatrix}$$

$$v_j = W_v x_j$$

$$= \begin{bmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \end{bmatrix} \begin{bmatrix} x_1^j \\ x_2^j \\ x_3^j \end{bmatrix} : 3 \times 1$$

$$y_i = \sum_j w_{ij} v_j$$

$$X : (B, T, L)$$

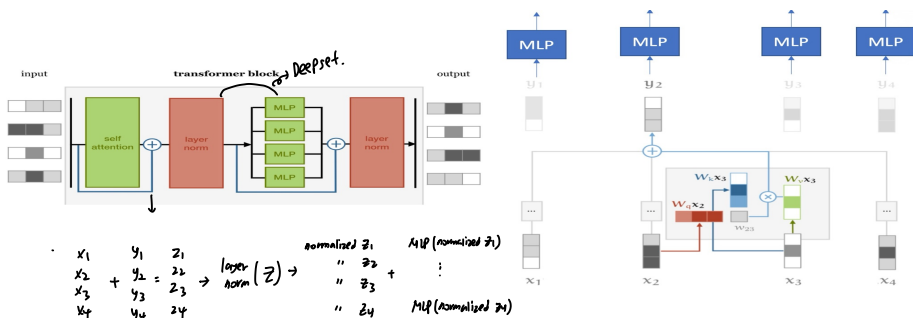
$$\text{SelfAttn}(C, h) \rightarrow B, T, H$$

$$\text{SelfAttn}(C, h) \rightarrow B, T, H$$

$$\begin{aligned} W &= Q \otimes K^T \\ &\rightarrow \text{Tensor Mat} \\ &: B \times L^2 \end{aligned}$$

Transformer

Self-attention layer \rightarrow Layer normalization \rightarrow Dense Layer



However, there are problems that order of sequence affect the result of computation, such as MLP.


: Let's encode each vector with position

Text, Signals, and Images are not sets.

: There is an inherent ordering on most domains we deal with

Yoda is a Jedi master! \neq a Jedi master Yoda is!

Position embedding




In "Attention is All you need" paper,

$$\vec{P}_k^{(i)} = f(k) := \begin{cases} \sin(w_k \cdot t) & \text{if } i = 2k \\ \cos(w_k \cdot t) & \text{if } i = 2k+1 \end{cases}$$

$$\text{where } w_k = \frac{1}{10000^{\frac{2k}{d}}}$$

$$\vec{P}_k = \begin{bmatrix} \sin(w_k \cdot t) \\ \cos(w_k \cdot t) \\ \vdots \\ \sin(w_{k+d/2} \cdot t) \\ \cos(w_{k+d/2} \cdot t) \end{bmatrix} \text{ dxl where } d \text{ be the encoding dimension } (d \approx 0) \text{ (embedding)}$$


Example)



Why trigonometric functions?

: Trigonometric functions like sin and cos naturally represent a pattern that the model can recognize as continuous.

Layer Normalization



Layer Normalization

Layer Norm: Calculate μ, σ^2 of Channel per each batch.

Example of layer norm:

Batches of 3 items:

Item ①

| |
|--------|
| 50.149 |
| 3314.8 |
| ... |
| ... |
| 8941.2 |
| 1444.7 |

μ_1
 σ_1^2

Item ②

| |
|--------|
| 1990.2 |
| 688.3 |
| ... |
| ... |
| 27.4 |
| 94.18 |

μ_2
 σ_2^2

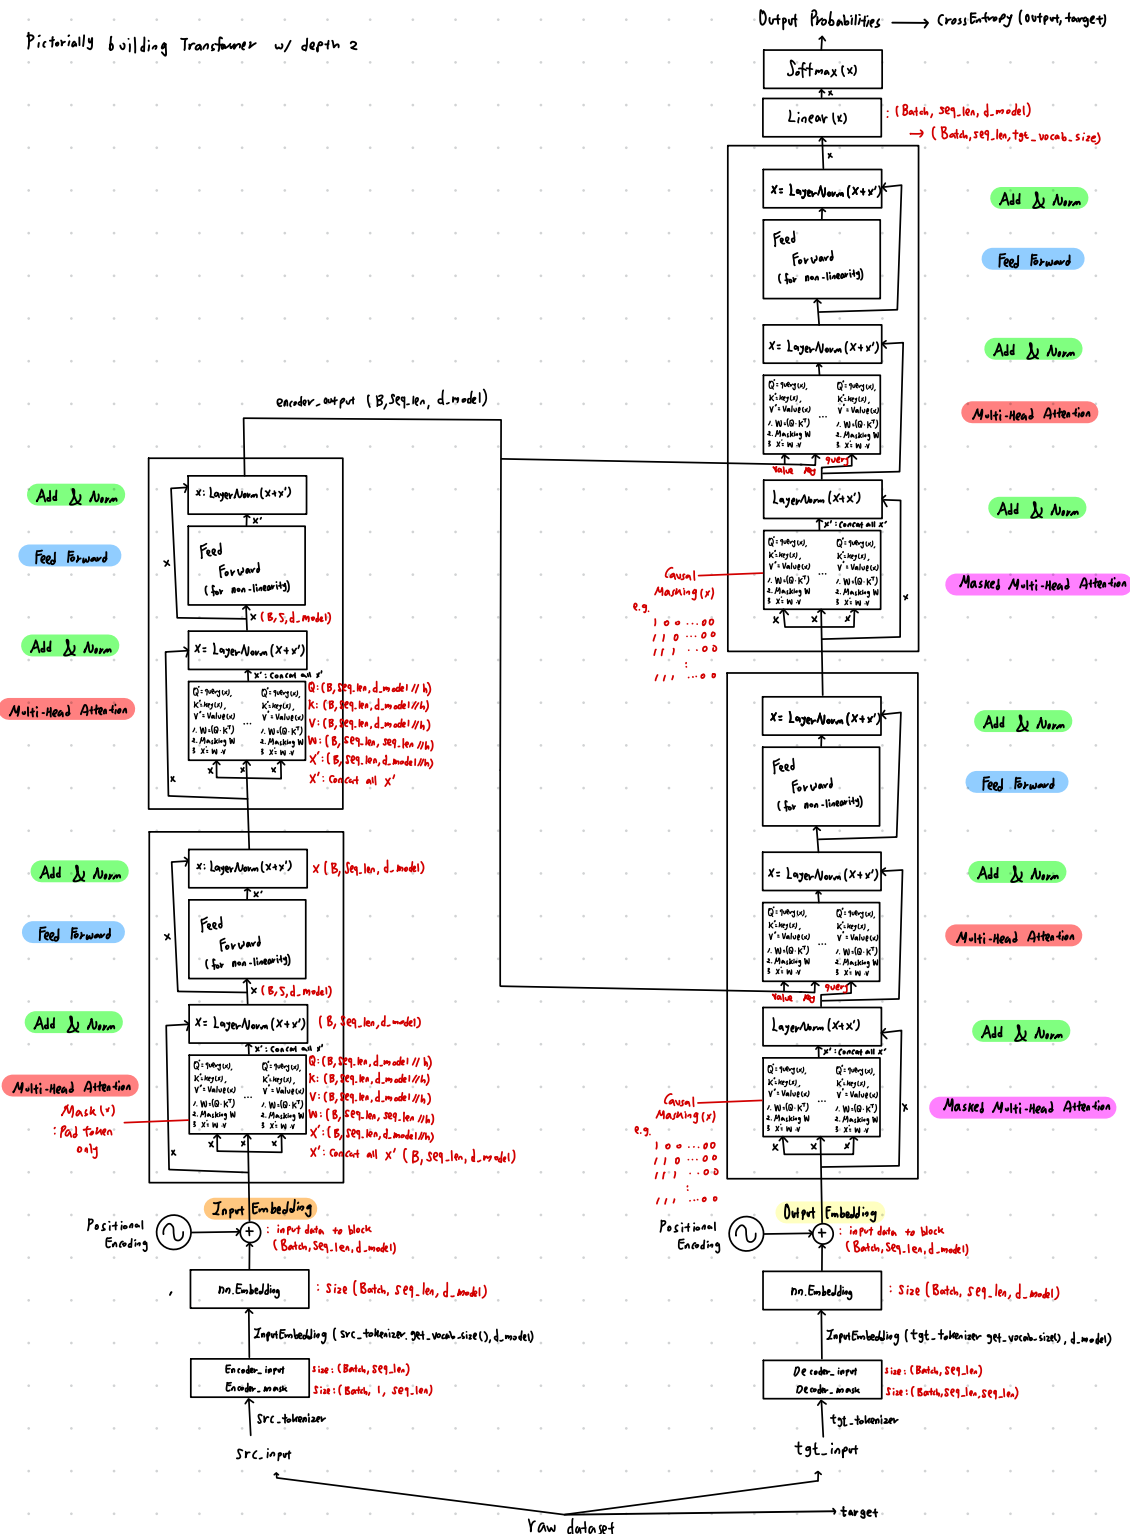
Item ③

| |
|---------|
| 182.7 |
| 174.878 |
| ... |
| ... |
| 10923.7 |
| 1004.88 |

μ_3
 σ_3^2

$$\Rightarrow \hat{x}_j = \frac{x_j - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

Pictorially building Transformer w/ depth 2



Vision Transformer (ViT)

In practice: take 224x224 input image,
divide into 14x14 grid of 16x16 pixel
patches (or 16x16 grid of 14x14 patches)

Each attention matrix has $14^4 = 38,416$
entries, takes 150 KB
(or 65,536 entries, takes 256 KB)

Output vectors

Exact same as
NLP Transformer!

Add positional
embedding: learned D-
dim vector per position

Linear projection to
D-dimensional vector

N input patches, each
of shape 3x16x16

Transformer

Linear projection
to C-dim vector
of predicted
class scores

Special extra input:
classification token
(D dims, learned)

Dosovitskiy et al, "An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale", ICLR 2021

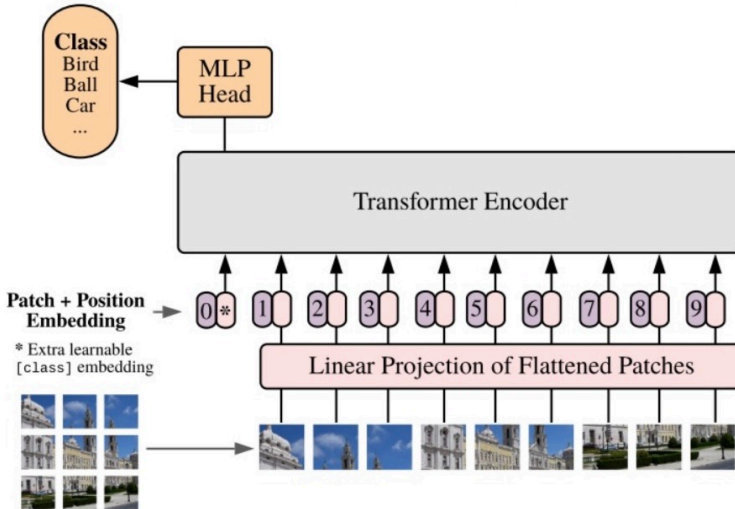
[Cat image](#) is free for commercial
use under a [Pixabay license](#)

2/27/24

Deep Learning - Lecture 11

21

Vision Transformer (ViT)



Transformer Encoder

